

GPU MODE Lecture 14: Practitioners Guide to Triton

NOTES

CUDA

Lecture #14 provides a practical introduction to writing and optimizing GPU kernels using Triton, featuring comparisons with CUDA and hands-on examples like tensor copying, image processing, and matrix multiplication.

AUTHOR

Christian Mills

PUBLISHED

November 15, 2024

This post is part of the following series:

- [GPU MODE Lecture Notes](#): My notes from the **GPU MODE** reading group lectures run by **Andreas Kopf** and **Mark Saroufim**.

- [Introduction](#)
- [Overview of the Talk](#)
- [Why and When to Use Triton](#)
- [How to Write Triton Kernels](#)
- [Practical Examples](#)
- [Benchmarking](#)
- [Auto-Tuning](#)
- [Conclusion and Resources](#)

Resource Links:

- **YouTube Recording:** [Lecture 14: Practitioners Guide to Triton](#)
- **Code:** [gpu-mode/lectures/lecture_014](#)

Introduction

- **Speaker:** [Umer Adil](#)
 - Former **management consultant** until October 2023.
 - Transitioned to **technical AI work** focusing on **open-source contributions**.
 - Contributed to projects like [LangChain](#) and [GPT Engineer](#).
 - Became a **maintainer** for GPT Engineer.
 - Currently contributing to **Hugging Face's** [diffusers](#).
 - Implemented features like **ControlNetXS**, **LoRAs**, etc.
 - Contact:
 - **GitHub:** [UmerHA](#)

- **Twitter:** [@UmerHAdil](#)
- **Support Umer:**
 - Independent open-source contributor.
 - **Ko-fi:** <https://ko-fi.com/umerha>

Overview of the Talk

- **Title:** A Practitioner's Guide to Triton
- **Agenda:**
 - **Why and When to Use Triton**
 - **How to Write Triton Kernels**
 - **Programming Model**
 - **Practical Examples:**
 - Copying a Tensor
 - Grayscale an Image
 - Fast Matrix-Matrix Multiplication Kernel
 - **Benchmarking and Auto-Tuning**
 - Performance Measurement
 - Kernel Optimization

Why and When to Use Triton

What is Triton?

- **Triton** is a language for programming **GPUs**.
 - More convenient than **CUDA**.
 - Allows writing **Python-like code** that compiles to **PTX** (Parallel Thread Execution).
 - PTX is the same intermediate representation used by CUDA.
- **Triton Compiler:**
 - **Optimizes code** by rearranging it for better performance without changing its meaning.
 - Targets the **same hardware** as CUDA.

Comparing Triton to CUDA

- **CUDA:**
 - Like a **high-end camera**.
 - Offers **thousands of knobs** for fine-grained control.
 - Achieves the **absolute best performance**.
 - **Harder** to write and debug.
- **Triton:**
 - Like a **high-end smartphone camera**.
 - **Easier** to use with fewer controls.
 - Provides **very good performance** with less effort.
 - **Easier** to write and debug.

`torch.compile()` vs. Triton

- `torch.compile()`:
 - Optimizes your **PyTorch code** but not the underlying **kernels**.
 - Changes your code to make the best use of existing GPU kernels.
 - Sometimes writes simple new kernels using **Triton**.
- **Triton**:
 - Allows writing **custom kernels** for performance-critical parts.
 - Offers more control over kernel behavior.

When to Use Triton

- **Optimization Steps**:
 1. Use `torch.compile()`:
 - Start by using `torch.compile()` to optimize your code.
 2. Adapt Your Code:
 - Rewrite code to be more suitable for `torch.compile()`.
 - E.g., eliminate **graph breaks** to enable **CUDA graphs**.
 3. Profile and Identify Bottlenecks:
 - Find slow parts of your code using profiling tools.
 - Write **custom Triton kernels** for these parts.
 4. Consider **CUDA**:
 - If still not fast enough, write custom **CUDA kernels**.
- **Note**: For maximum performance from the start, you may choose **CUDA** directly.

Rough Edges in Triton

- **New-ish Project**:
 - Contains **rough edges**; code may not behave as expected.
 - Expected to become more polished over time.
- **Recommendation**:
 - **Debugging** is important; use “simulator mode” when possible.
 - Be aware of limitations on older GPUs or with certain operations.

How to Write Triton Kernels

Debugging Triton Kernels

- **Simulator Mode**:
 - Set environment variable `TRITON_INTERPRET='1'`.
 - Enables debugging by running kernels on the **CPU**.
- **Advantages**:
 - **Debug** and **print** variables like in CPU programs.
 - Easier to **set breakpoints** and inspect program flow.
- **Utility Functions**:

```

import os

os.environ['TRITON_INTERPRET'] = '1' # needs to be set *before* triton is imported

import triton
import triton.language as tl
from IPython.core.debugger import import set_trace

def test_pid_conds(conds, pid_0=[0], pid_1=[0], pid_2=[0]):
    """
    Test if conditions on program IDs (PIDs) are fulfilled.

    Args:
        conds (str): String containing conditions to check. Multiple conditions are separated by spaces. Each condition consists of an operator and a number.
        pid_0 (list): First program ID value in a single-element list. Default: [0]
        pid_1 (list): Second program ID value in a single-element list. Default: [0]
        pid_2 (list): Third program ID value in a single-element list. Default: [0]

    Examples:
        '=0'      -> Checks if pid_0 equals 0
        ',>1'     -> Checks if pid_1 is greater than 1
        '>1,=0'   -> Checks if pid_0 > 1 AND pid_1 = 0

    Returns:
        bool: True if all conditions are met, False otherwise
    """
    # Extract PID values from lists
    pids = pid_0[0], pid_1[0], pid_2[0]

    # Remove spaces and split conditions by comma
    conds = conds.replace(' ', '').split(',')

    # Check each condition against corresponding PID
    for i, (cond, pid) in enumerate(zip(conds, pids)):
        if cond == '': continue # Skip empty conditions

        # Split condition into operator and threshold value
        op, threshold = cond[0], int(cond[1:])

        # Validate operator
        valid_ops = ['<', '>', '>=', '<=', '=', '!=']
        if op not in valid_ops:
            raise ValueError(f"Rules may only use these ops: {valid_ops}. Invalid rule: {cond}")

        # Convert '=' to '==' for Python evaluation
        op = '==' if op == '=' else op

        # Evaluate condition
        if not eval(f'{pid} {op} {threshold}'):

```

```
return False
```

```
return True
```

- `check_tensors_gpu_ready`:

```
def check_tensors_gpu_ready(*tensors):
    """
    Verify that all input tensors are contiguous and on GPU.

    Args:
        *tensors: Variable number of PyTorch tensors to check

    Raises:
        AssertionError: If any tensor is not contiguous or not on GPU
    """
    for t in tensors:
        assert t.is_contiguous(), "A tensor is not contiguous"
        # Skip GPU check if in simulator mode
        if not os.environ.get('TRITON_INTERPRET') == '1':
            assert t.is_cuda, "A tensor is not on cuda"
```

- Ensure data is ready for GPU execution.
 1. Assert all tensors are contiguous in memory
 2. Assert all tensors are on GPU, if not simulating

- `print_if`:

```
def print_if(txt, conds, pid_0=[0], pid_1=[0], pid_2=[0]):
    """
    Print text if specified PID conditions are met.
    Useful for debugging specific threads in GPU kernels.

    Args:
        txt (str): Text to print
        conds (str): Conditions string (same format as test_pid_conds)
        pid_0, pid_1, pid_2 (list): Program ID values to check
    """
    if test_pid_conds(conds, pid_0, pid_1, pid_2):
        print(txt)
```

- Output variable values for debugging, depending on conditions on PIDs

- `breakpoint_if`:

```
def breakpoint_if(conds, pid_0=[0], pid_1=[0], pid_2=[0]):
    """
    Set a breakpoint if specified PID conditions are met.
    Useful for debugging specific threads in GPU kernels.
```

```

Args:
    conds (str): Conditions string (same format as test_pid_conds)
    pid_0, pid_1, pid_2 (list): Program ID values to check
    """
    if test_pid_conds(conds, pid_0, pid_1, pid_2):
        set_trace()

```

- Pause execution at specific points, depending on conditions on PIDs

Programming Model

CUDA vs. Triton

- **CUDA:**
 - **Two-tiered Decomposition:**
 - **Blocks:** Groups of threads.
 - **Threads:** Operate on **scalar values**.
 - Threads within a block share the same **Streaming Multiprocessor (SM)** and **shared memory**.
- **Triton:**
 - **One-tiered Decomposition:**
 - **Programs** (equivalent to blocks in CUDA).
 - Operates on **vectors** instead of scalars.
 - **Vectorized Operations:**
 - All operations (loading, computing, storing) are performed on vectors.
 - No explicit management of threads or shared memory.

Example: Adding Two Vectors

- **Task:**
 - Add vectors **X** and **Y** of size **8** to produce **Z**.
- **CUDA Approach:**

```
def add_cuda_k(x, y, z, n, bs):
```

```
    """
```

A simplified Python representation of a CUDA kernel that adds two vectors element-wise. This function demonstrates how parallel processing works in CUDA, where multiple threads process different parts of the data simultaneously.

Parameters:

```

    x (array): First input vector to be added
    y (array): Second input vector to be added
    z (array): Output vector where results will be stored
    n (int): Total size of the input vectors
    bs (int): Block size – number of threads per block
               (determines how many elements each block processes)

```

Example:

```

    If n = 8 and bs = 4:

```

```

- We'll have 2 blocks (block_id: 0,1)
- Each block has 4 threads (thread_id: 0,1,2,3)
- Each thread processes one element
  block 0 handles indices 0-3, block 1 handles indices 4-7
"""
# In CUDA, these values would be automatically set based on the GPU's
# thread and block configuration. Here they're placeholders.
block_id = ... # Block ID (e.g., 0 or 1 if we have 2 blocks)
thread_id = ... # Thread ID within the block (e.g., 0,1,2,3 if block size is 4)

# Calculate which element this specific thread should process
# Example: If block_id=1, bs=4, thread_id=2:
# offs = 1 * 4 + 2 = 6 (this thread processes the 7th element)
offs = block_id * bs + thread_id

# Only process if we're within the vector bounds
# This check is necessary because the number of threads might be more
# than the actual data we need to process
if offs < n:
    # Each thread reads its assigned values from the input vectors
    x_value = x[offs] # Get value from first input vector
    y_value = y[offs] # Get value from second input vector

    # Perform the addition operation
    z_value = x_value + y_value

    # Store the result in the output vector
    z[offs] = z_value

# Note: In actual CUDA programming, all variables above are scalars
# (single values, not arrays). Each thread works with just one element,
# but many threads run in parallel to process the entire array quickly.

```

- Use **2 blocks** with **4 threads** each (block size of 4).
- Each thread computes on a **scalar value**.

block_id	thread_id	offs	offs < n	x_value	y_value	z_value
0	0	0	True	1	0	1
0	1	1	True	2	1	3
0	2	2	True	3	0	3
0	3	3	True	4	1	5
1	0	4	True	5	0	5
1	1	5	True	6	1	7
1	2	6	False	undefined	undefined	undefined
1	3	7	False	undefined	undefined	undefined

- Triton Approach:

```
def add_triton_k(x, y, z, n, bs):
```

```
    """
```

A representation of a Triton kernel that adds two vectors element by element. Unlike CUDA, Triton operates on vectors (groups of elements) rather than individuals which can lead to more efficient code execution.

Parameters:

x (array): First input vector to be added
y (array): Second input vector to be added
z (array): Output vector where results will be stored
n (int): Total size of the input vectors
bs (int): Block size – number of elements to process in each block
(This determines the size of the vectors we operate on)

Key Differences from CUDA:

- Triton processes vectors (multiple elements at once) instead of single values
- Operations are vectorized, meaning they work on entire arrays simultaneously
- No explicit thread_id needed as Triton handles multiple elements per block

```
    """
```

```
# Get the block ID for this kernel instance
```

```
# tl.program_id(0) is Triton's way of identifying which block we're processing
```

```
block_id = tl.program_id(0) # Example: 0 or 1 if processing in two blocks
```

```
# Create a vector of offsets for this block
```

```
# tl.arange(0, bs) creates a vector like [0, 1, 2, ..., bs-1]
```

```
# For example, if block_id=1 and bs=4:
```

```
# offs = 4 + [0,1,2,3] = [4,5,6,7]
```

```
offs = block_id * bs + tl.arange(0, bs) # Vector of indices to process
```

```
# Create a mask for valid elements
```

```
# This returns a vector of boolean values
```

```
# Example: if n=6 and offs=[4,5,6,7], mask=[True,True,False,False]
```

```
mask = offs < n # Vector of bools showing which elements are valid
```

```
# Load multiple elements at once from input vectors
```

```
# x[offs] loads multiple values in parallel
```

```
# Example: if offs=[4,5,6,7], this loads four elements at once
```

```
x_values = x[offs] # Load vector of values from first input
```

```
y_values = y[offs] # Load vector of values from second input
```

```
# Perform vectorized addition
```

```
# This adds entire vectors element-wise in one operation
```

```
# Example: [1,2,3,4] + [5,6,7,8] = [6,8,10,12]
```

```
z_value = x_value + y_value # Add vectors element-wise
```

```
# Store results back to memory
```

```
# Writes multiple elements at once
```

```
# The mask ensures we only write valid results
```



```
z[offs] = z_value # Store vector of results
```

```
# Note: All operations above work on vectors (multiple elements at once)
# This is more efficient than CUDA's scalar operations because:
# 1. Fewer memory transactions are needed
# 2. Vector operations can utilize SIMD instructions
# 3. Less overhead from individual thread management
```

- Use **2 programs** (no threads).
- Each program operates on a **vector** of size **4**.
- **Offsets** and **masks** are vectors.

block_id	offs	offs < n	x_value	y_value	z_value
0	[0, 1, 2, 3]	[True, True, True, True]	[1, 2, 3, 4]	[0, 1, 0, 1]	[1, 3, 3, 5]
1	[4, 5, 6, 7]	[True, True, False, False]	[5, 6, 'undefined', 'undefined']	[0, 1, 'undefined', 'undefined']	[5, 7, 'undefined', 'undefined']

Jargon

- **Program:**
 - A kernel instance processing a block of data.
- **PID (Program ID):**
 - Equivalent to **Block ID** in CUDA.
- **Vectorized Operations:**
 - Simultaneous operations on multiple data points.

Practical Examples

```
import torch
import triton
import triton.language as tl
```

```
def cdiv(n, d):
    """
    Compute ceiling division between two numbers.
    Args:
        n: Numerator
        d: Denominator
    Returns:
        Ceiling division result
    """
    return (n + d - 1) // d
```

```
x = torch.tensor([1,2,3,4,5,6])
y = torch.tensor([0,1,0,1,0,1])
```

```
x, y, x+y
```

```
(tensor([1, 2, 3, 4, 5, 6]),
 tensor([0, 1, 0, 1, 0, 1]),
 tensor([1, 3, 3, 5, 5, 7]))
```

Example 1: Copying a Tensor

```
def copy(x, bs, kernel_fn):
    """
    Launch a Triton kernel to copy data from one GPU tensor to another.

    Args:
        x: Input tensor to copy from
        bs: Block size – number of elements processed per GPU thread block
        kernel_fn: Triton kernel function to execute

    Returns:
        z: New tensor containing copied data
    """
    # Create output tensor with same properties as input
    z = torch.zeros_like(x)

    # Verify tensors are GPU-ready
    check_tensors_gpu_ready(x, z)

    # Calculate grid dimensions for GPU execution
    n = x.numel() # Total number of elements
    n_blocks = cdiv(n, bs) # Number of thread blocks needed
    grid = (n_blocks,) # 1D grid configuration

    # Launch kernel on GPU
    kernel_fn[grid](x, z, n, bs)
    return z
```

Objective

- Copy tensor **X** of shape **N** to tensor **Z**.

Steps

1. Define Kernel:

- Use `@triton.jit` decorator.
- Function arguments are pointers to tensors and size parameters.

2. Calculate Offsets:

- Compute offsets using **PID** and **block size**.
- `offsets = pid * block_size + tl.arange(0, block_size)`

3. Create Mask:

- Prevent out-of-bounds access.
- `mask = offsets < N`

4. Load and Store Data:

- Load data from **X**: `x = tl.load(X + offsets, mask=mask)`
- Store data to **Z**: `tl.store(Z + offsets, x, mask=mask)`

5. Launch Kernel:

- Determine grid size: `grid = (num_blocks,)`
- Call kernel with `grid` and `block_size`.

Debugging

- **Intentional Bug:** Incorrect offset calculation

```
# Basic kernel with incorrect offset calculation
@triton.jit # This decorator converts the Python function into GPU code
def copy_k(x_ptr, z_ptr, n, bs: tl.constexpr):
    """
    Initial version of copy kernel – demonstrates common mistake.

    Important Notes:
    - The @triton.jit decorator transforms this Python function into GPU code
    - Only a limited set of operations are allowed inside GPU kernels:
      * Basic arithmetic and logic operations are allowed
      * Python print() and debugging tools like breakpoints are NOT allowed
      * Use specialized Triton functions for GPU operations

    Args:
      x_ptr: Pointer to input tensor data (Triton automatically converts tensor to
      z_ptr: Pointer to output tensor data
      n: Total number of elements
      bs: Block size (marked as compile-time constant with tl.constexpr)

    Note: This version has a bug – it processes the same elements in each block!
    """
    pid = tl.program_id(0) # Get current block ID
    offs = tl.arange(0, bs) # Creates offsets [0, 1, ..., bs-1]
    mask = offs < n # Prevent out-of-bounds access
    x = tl.load(x_ptr + offs, mask) # Load input values
    tl.store(z_ptr + offs, x, mask) # Store to output

    print_if(f'pid = {pid} | offs = {offs}, mask = {mask}, x = {x}', '')
```

```
z = copy(x, bs=2, kernel_fn=copy_k)
```

```
pid = [0] | offs = [0 1], mask = [ True  True], x = [1 2]
pid = [1] | offs = [0 1], mask = [ True  True], x = [1 2]
pid = [2] | offs = [0 1], mask = [ True  True], x = [1 2]
```

```
z
```

```
tensor([1, 2, 0, 0, 0, 0])
```

- Incorrectly calculating offsets without considering **PID**.
- Only the first block of data is copied.
- **Intentional Bug:** Incorrect stride calculation

```
# Incorrect stride calculation
@triton.jit
def copy_k(x_ptr, z_ptr, n, bs: tl.constexpr):
    """
    Second version – demonstrates another common mistake.

    Key Concepts:
    - When we pass a torch tensor to the kernel, Triton automatically converts it
      to a pointer to its first element (that's why we receive x_ptr, not x)
    - GPU kernels run in parallel across many blocks, so correct memory access
      patterns are crucial

    Note: This version incorrectly uses 'n' instead of 'bs' for stride calculation,
    causing blocks to process wrong sections of memory.
    """
    pid = tl.program_id(0)
    offs = pid * n + tl.arange(0, bs) # Wrong! Stride should use 'bs', not 'n'
    mask = offs < n
    x = tl.load(x_ptr + offs, mask)
    tl.store(z_ptr + offs, x, mask)

    print_if(f'pid = {pid} | offs = {offs}, mask = {mask}, x = {x}', '')
```

```
z = copy(x, bs=2, kernel_fn=copy_k)
```

```
pid = [0] | offs = [0 1], mask = [ True  True], x = [1 2]
pid = [1] | offs = [6 7], mask = [False False], x = [0 0]
pid = [2] | offs = [12 13], mask = [False False], x = [0 0]
```

- **Solution:**

```
# Correct implementation
@triton.jit
def copy_k(x_ptr, z_ptr, n, bs: tl.constexpr):
    """
    Final correct version of the copy kernel.

    GPU Kernel Rules and Concepts:
    1. The @triton.jit decorator converts this Python function into GPU code
    2. Inside GPU kernels:
       - You can't use regular Python print() or debuggers
       - You must use special Triton functions (tl.*) for operations
       - Tensor inputs are automatically converted to memory pointers
```

```

3. Each block processes a different chunk of data in parallel:
    - Block 0 processes elements [0:bs]
    - Block 1 processes elements [bs:2*bs]
    - Block 2 processes elements [2*bs:3*bs]
    ....
pid = tl.program_id(0) # Get current block ID
offs = pid * bs + tl.arange(0, bs) # Calculate correct offsets for this block
mask = offs < n # Prevent out-of-bounds access
x = tl.load(x_ptr + offs, mask) # Load input values
tl.store(z_ptr + offs, x, mask) # Store to output

print_if(f'pid = {pid} | offs = {offs}, mask = {mask}, x = {x}', '')

```

```
z = copy(x, bs=2, kernel_fn=copy_k)
```

```

pid = [0] | offs = [0 1], mask = [ True  True], x = [1 2]
pid = [1] | offs = [2 3], mask = [ True  True], x = [3 4]
pid = [2] | offs = [4 5], mask = [ True  True], x = [5 6]

```

- Adjust offsets to include `pid * block_size`.

Key Takeaways

- **Offset Calculation** is crucial.
- Use **masks** to handle data boundaries.
- **Debugging** is facilitated by simulator mode.

Example 2: Grayscaleing an Image

Note

Umer mentioned needing to restart the notebook kernel before running this example because:

1. *torchvision can't be imported, probably due to a circular dependency. -> I currently don't know why, need to dig deeper.*
2. *the simulated triton kernel below fails, because a float can't be multiplied to a uint vector -> Works on GPU w/o simulation, so seems to be a `TRITON_INTERPRET` bug.*

However, the underlying issues seem to have been resolved in more recent updates.

```

# Import required libraries
import os
import matplotlib.pyplot as plt
from urllib.request import urlretrieve # For downloading files from URLs
from pathlib import Path
import torch
from torch import tensor
import torchvision as tv

```

```
import torchvision.transforms.functional as tvf
from torchvision import io
import triton # GPU acceleration library
import triton.language as tl
```

```
# Define image URL and download if not already present
url = 'https://upload.wikimedia.org/wikipedia/commons/thumb/4/43/Cute_dog.jpg/1600px-Cute'
path_img = Path('puppy.jpg')
if not path_img.exists():
    urlretrieve(url, path_img)
```

```
# Read the image using torchvision
img = io.read_image('puppy.jpg')
print(f"Image shape (channels, height, width): {img.shape}")
img[:2,:3,:4]
```

```
torch.Size([3, 1066, 1600])
```

```
tensor([[[[117, 119, 117, 113],
          [119, 129, 129, 113],
          [130, 126, 122, 115]],

        [[ 83,  85,  85,  80],
          [ 85,  97,  97,  82],
          [ 98,  93,  89,  83]]], dtype=torch.uint8)
```

```
def show_img(x, figsize=(4,3), **kwargs):
    """
    Display an image using matplotlib

    Args:
        x: Image tensor
        figsize: Figure size in inches (width, height)
        **kwargs: Additional arguments passed to plt.imshow()
    """
    plt.figure(figsize=figsize)
    plt.axis('off')
    # Convert from CHW (channels, height, width) to HWC format if needed
    if len(x.shape) == 3:
        x = x.permute(1, 2, 0)
    plt.imshow(x.cpu(), **kwargs)
```

Objective

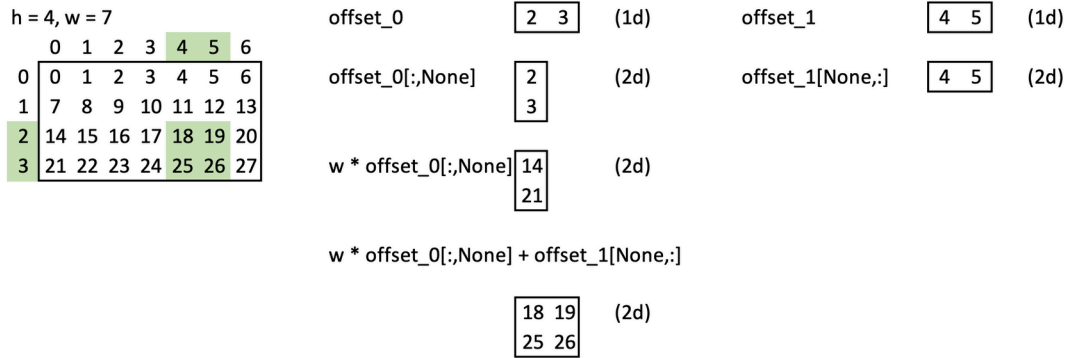
- Convert a color image to grayscale using a Triton kernel.

Steps

1. Load Image Data:

- Use an image (e.g., a puppy image) as input.

2. Calculate 2D Offsets:



- Compute row and column offsets.
- Use broadcasting to create a grid of offsets.

3. Create Masks:

- Handle image boundaries to avoid out-of-bounds access.

4. Load Color Channels:

- Load **R**, **G**, and **B** values using offsets.

5. Compute Grayscale Values:

- Apply formula: $\text{grayscale} = 0.2989 \cdot R + 0.5870 \cdot G + 0.1140 \cdot B$

6. Store Grayscale Data:

- Write the grayscale values back to the output tensor.

Implementation

```
@triton.jit
def rgb2grey_k(x_ptr, out_ptr, h, w, bs0: tl.constexpr, bs1: tl.constexpr):
    """
    GPU kernel for converting RGB image to grayscale

    Args:
        x_ptr: Pointer to input RGB image data
        out_ptr: Pointer to output grayscale image data
        h: Image height
        w: Image width
        bs0: Block size for height dimension
        bs1: Block size for width dimension
    """
    # Get program IDs for parallel processing
    pid_0 = tl.program_id(0) # Block ID in height dimension
    pid_1 = tl.program_id(1) # Block ID in width dimension

    # Calculate offsets for this block
    offs_0 = pid_0 * bs0 + tl.arange(0, bs0) # Offsets in height dimension
    offs_1 = pid_1 * bs1 + tl.arange(0, bs1) # Offsets in width dimension
```

```

# Calculate 2D offset matrix
offs = w * offs_0[:,None] + offs_1[None, :]

# Create masks to handle image boundaries
mask_0 = offs_0 < h
mask_1 = offs_1 < w
mask = mask_0[:,None] & mask_1[None,:]

# Load RGB channels
r = tl.load(x_ptr + 0*h*w + offs, mask=mask)
g = tl.load(x_ptr + 1*h*w + offs, mask=mask)
b = tl.load(x_ptr + 2*h*w + offs, mask=mask)

# Convert to grayscale using standard weights
# These weights represent human perception of color:
# Red: 29.89%, Green: 58.70%, Blue: 11.40%
out = 0.2989*r + 0.5870*g + 0.1140*b

# Store the result
tl.store(out_ptr + offs, out, mask=mask)

```

Notes

- **Vectorized Operations** simplify processing of 2D data.
- **Masks** ensure safe memory access.
- **GPU Compatibility:**
 - Some operations may not work in simulator mode or on older GPUs.

Execution

- **Kernel Launch:**

```

def rgb2grey(x, bs):
    """
    Convert RGB image to grayscale using GPU acceleration

    Args:
        x: Input RGB image tensor (channels, height, width)
        bs: Tuple of block sizes (height, width) for GPU processing

    Returns:
        Grayscale image tensor (height, width)
    """
    c, h, w = x.shape
    # Create output tensor
    out = torch.empty((h,w), dtype=x.dtype, device=x.device)

    # Define processing grid based on block sizes
    grid = lambda meta: (cdiv(h, meta['bs0']), cdiv(w, meta['bs1']))

```



```
# Launch GPU kernel
rgb2grey_k[grid](x, out, h, w, bs0=bs[0], bs1=bs[1])
return out.view(h,w)
```

- Define grid dimensions based on image size.

- **Result:**

```
# Resize image to a smaller size for faster processing
img = tvf.resize(img, 150, antialias=True)
ch, h, w = img.shape # Get channels, height, and width
ch,h,w,h*w
```

(3, 150, 225, 33750)

```
show_img(img)
```



```
# Convert image to grayscale and display
grey_img = rgb2grey(img.to('cuda'), bs=(32, 32)).to('cpu')
show_img(grey_img, cmap='gray')
```



- Successfully converted grayscale image.

Example 3: Matrix Multiplication

Note

Had to restart the notebook kernel to produce expected results for this example.

```
import os
import torch
import triton
import triton.language as tl
```

Objective

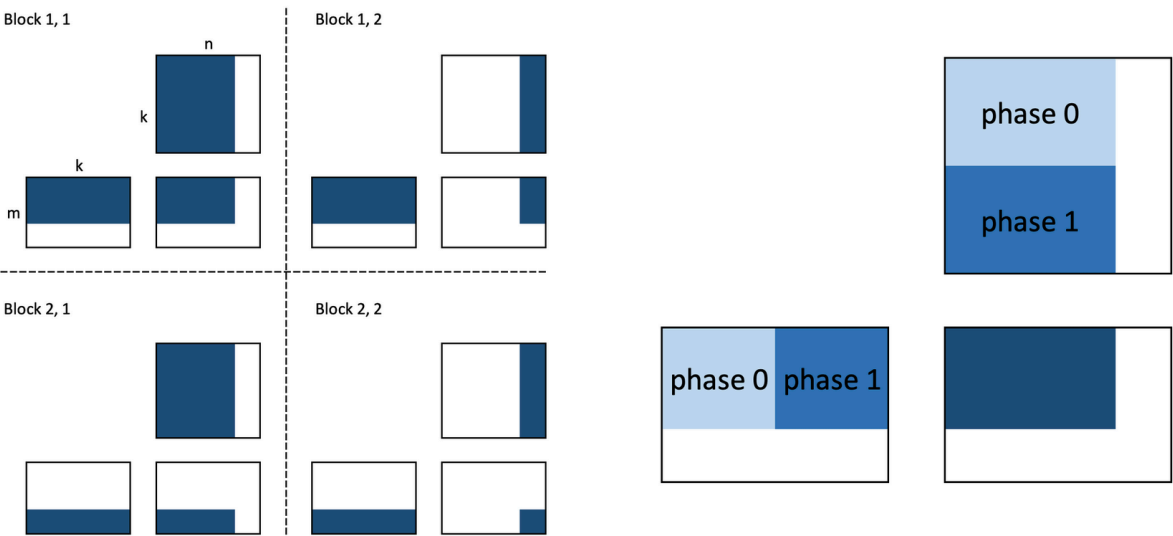
- Implement an efficient matrix multiplication kernel.

Decomposition Strategy

- **Matrices:**
 - **A:** Size $M \times K$
 - **B:** Size $K \times N$
 - **C:** Result $M \times N$
- **Splitting:**

Splits along m- and n-dimensions are represented by blocks

Split along k-dimension is represented by ‘phases’, which are done in the same block



- Split **C** along **M** and **N** dimensions.
- Map splits to programs (blocks).
- Further split along **K** dimension (phases).

Naive Matrix Multiplication

1. **Define Kernel:**
 - Use `@triton.jit` decorator.

2. Calculate Offsets:

- Compute offsets for **M** and **N** axes.

3. Initialize Accumulator:

- Set to zero before accumulation.

4. Loop Over K Dimension:

- For each phase:
 - Load chunks of **A** and **B**.
 - Multiply and accumulate.

5. Store Result:

- Write the computed block to **C**.

Implementation

• Helper Functions:

```
# ===== Helper Functions for Computing Memory Offsets and Masks =====

@triton.jit
def get_1d_offset(size, n_prev_chunks):
    """
    Calculate 1D memory offsets for a given chunk size and position.

    Args:
        size: Size of the current chunk
        n_prev_chunks: Number of previous chunks (used for position)

    Returns:
        Array of offsets for the current chunk
    """
    return n_prev_chunks * size + tl.arange(0, size)

@triton.jit
def get_2d_offset(off0, off1, stride0, stride1=1):
    """
    Calculate 2D memory offsets for matrix operations.

    Args:
        off0, off1: Offsets in first and second dimensions
        stride0, stride1: Stride values for memory layout

    Returns:
        2D array of memory offsets
    """
    return tl.expand_dims(off0, 1)*stride0 + tl.expand_dims(off1, 0)*stride1

@triton.jit
def get_1d_mask(off, max):
    """
    Create a mask for boundary checking in 1D.
```

Args:

- offs: Current offsets
- max: Maximum valid offset

Returns:

- Boolean mask indicating valid positions

```

"""
    return offs < max

```

```

@triton.jit
def get_2d_mask(offs_0, offs_1, max_0, max_1):
    """
    Create a mask for boundary checking in 2D.

    Args:
        offs_0, offs_1: Current offsets in both dimensions
        max_0, max_1: Maximum valid offsets

    Returns:
        Boolean mask indicating valid positions in 2D
    """
    return (tl.expand_dims(offs_0, 1) < max_0) & (tl.expand_dims(offs_1, 0) < max_1)

```

- **Matrix Multiplication Kernel:**

```

@triton.jit
def naive_matmul_k(
    a_ptr, b_ptr, c_ptr, # Pointers to input/output matrices
    m, n, k,             # Matrix dimensions: A(m×k), B(k×n), C(m×n)
    stride_am, stride_ak, # Memory strides for matrix A
    stride_bk, stride_bn, # Memory strides for matrix B
    stride_cm, stride_cn, # Memory strides for output matrix C
    bm: tl.constexpr,     # Block size for M dimension
    bn: tl.constexpr,     # Block size for N dimension
    bk: tl.constexpr      # Block size for K dimension
):
    """
    Compute matrix multiplication C = A × B using block-wise operations.

    This kernel implements a basic matrix multiplication by:
    1. Breaking the computation into blocks
    2. Loading blocks into shared memory
    3. Computing partial results
    4. Storing the results back to global memory

    Args:
        a_ptr, b_ptr: Input matrix pointers
        c_ptr: Output matrix pointer
        m, n, k: Matrix dimensions
        stride_*: Memory strides for each matrix
    """

```

```

    bm, bn, bk: Block sizes for tiled computation
    """
    # Get program ID for the current thread block
    pid_m, pid_n = tl.program_id(0), tl.program_id(1)

    # Calculate offsets for the current block
    rm = get_1d_offset(size=bm, n_prev_chunks=pid_m) # Offset in M dimension
    rn = get_1d_offset(size=bn, n_prev_chunks=pid_n) # Offset in N dimension
    rk = get_1d_offset(size=bk, n_prev_chunks=0)      # Initial offset in K dimension

    # Calculate memory offsets for input matrices
    offs_a = a_ptr + get_2d_offset(rm, rk, stride_am, stride_ak)
    offs_b = b_ptr + get_2d_offset(rk, rn, stride_bk, stride_bn)

    # Initialize accumulator for partial results
    # Note: allow_tf32 must be set to False for older GPUs
    acc = tl.zeros((bm, bn), dtype=tl.float32)

    # Main computation loop - iterate over K dimension
    for _ in range(0, k, bk):
        # Load blocks from input matrices
        a = tl.load(offs_a) # Load block from matrix A
        b = tl.load(offs_b) # Load block from matrix B

        # Compute partial matrix multiplication for current block
        acc += tl.dot(a, b, allow_tf32=False)

        # Update offsets for next iteration
        offs_a += bk * stride_ak
        offs_b += bk * stride_bk

    # Calculate output memory location and mask for boundary conditions
    c = c_ptr + get_2d_offset(rm, rn, stride_cm, stride_cn)
    mask = get_2d_mask(rm, rn, m, n)

    # Store the result
    tl.store(c, acc, mask=mask)

```

```

def matmul(a, b, matmul_k_fn, bs=16, group_sz=None):
    """
    High-level matrix multiplication function that handles kernel launch.

    Args:
        a, b: Input matrices
        matmul_k_fn: Triton kernel function to use
        bs: Block size for tiled computation
        group_sz: Group size for advanced implementations

    Returns:
        Result of matrix multiplication
    """

```

```

"""
# Verify matrix dimensions are compatible
assert a.shape[1] == b.shape[0], "matrix dims not compatible for matmul"
check_tensors_gpu_ready(a, b)

# Get matrix dimensions
(m, k), (_, n) = a.shape, b.shape

# Initialize output matrix
c = torch.empty((m, n), device=a.device, dtype=torch.float16)

# Calculate grid dimensions for kernel launch
grid = lambda meta: (triton.cdiv(m, meta['bm']), triton.cdiv(n, meta['bn']))

# Handle optional group size parameter
group_sz = {} if group_sz is None else {"group_sz": group_sz}

# Launch kernel
matmul_k_fn[grid](
    a, b, c,                # Input/output matrices
    m, n, k,                # Matrix dimensions
    a.stride(0), a.stride(1), # Strides for matrix A
    b.stride(0), b.stride(1), # Strides for matrix B
    c.stride(0), c.stride(1), # Strides for output matrix
    bm=bs, bn=bs, bk=bs,    # Block sizes
    **group_sz
)

return c

```

- **Results:**

```

# Create a simplified interface using partial application
naive_matmul = partial(matmul, matmul_k_fn=naive_matmul_k)

```

```

# Small example
a = torch.ones((3, 4), dtype=torch.float32, device='cuda')
b = torch.ones((4, 5), dtype=torch.float32, device='cuda')
naive_matmul(a, b)

```

```

tensor([[4., 4., 4., 4., 4.],
        [4., 4., 4., 4., 4.],
        [4., 4., 4., 4., 4.]], device='cuda:0', dtype=torch.float16)

```

```

# Larger example with verification
torch.manual_seed(0)
a = torch.randn((512, 512), device='cuda', dtype=torch.float16)
b = torch.randn((512, 512), device='cuda', dtype=torch.float16)

```

```

# Compare Triton implementation with PyTorch

```

```

triton_output = naive_matmul(a, b)
torch_output = torch.matmul(a, b)

# Verify results match within tolerance
if torch.allclose(triton_output, torch_output, atol=5e-2, rtol=0):
    print("✅ Triton and Torch match")
else:
    print("❌ Triton and Torch differ")

```

✅ Triton and Torch match

Example 4: Faster Matrix Multiplication

Swizzling for Cache Optimization

- **Goal:**
 - Improve **L2 cache** utilization.
- **Swizzling:**
 - Reorder program execution to process blocks that share data closer in time.

With **naive ordering**

To compute 9 blocks of C

we need to load **9 blocks** from A

and **81 blocks** from B,

so **90 in total**.

0	1	2	3	4	5	6	7	8

=

x

With **grouped ordering**

To compute 9 blocks of C

we need to load **27 blocks** from A

and **27 blocks** from B,

so **54 in total**.

0	1	2						
3	4	5						
6	7	8						

=

x


```

def process_item(id):
    print(f"I'm processing item {id}")

# Demonstrate normal sequential processing
print("Sequential processing:")
for i in range(5):
    process_item(i)

```

```

I'm processing item 0
I'm processing item 1
I'm processing item 2
I'm processing item 3
I'm processing item 4

```

```

def change_id(old_id):
    return 5-old_id

# Demonstrate reordered processing
print("\nReordered processing:")
for i in range(5):
    process_item(change_id(i))


```

```

I'm processing item 5
I'm processing item 4
I'm processing item 3
I'm processing item 2
I'm processing item 1

```

- Implementation:



```

# ===== Memory Access Pattern Optimization via Swizzling =====
@triton.jit
def swizzle_k(x_ptr, z_ptr, group_sz: tl.constexpr):
    """
    Demonstrates memory access pattern optimization using swizzling.
    Swizzling reorders thread blocks to improve memory locality and cache utilization

    Args:
        x_ptr: Input tensor pointer
        z_ptr: Output tensor pointer
        group_sz: Size of thread block groups for swizzling
    """
    # Get current thread block coordinates and grid dimensions
    pid_m, pid_n = tl.program_id(0), tl.program_id(1)
    num_pid_m, num_pid_n = tl.num_programs(0), tl.num_programs(1)

    # Apply 2D swizzling to reorder thread blocks
    pid_m_, pid_n_ = tl.swizzle2d(pid_m, pid_n, num_pid_m, num_pid_n, group_sz)

```



```

# Calculate memory offsets for original ordering
offs_m = get_1d_offset(1, n_prev_chunks=pid_m)
offs_n = get_1d_offset(1, n_prev_chunks=pid_n)
offs = get_2d_offset(offs_m, offs_n, stride_0=num_pid_n)
mask = get_2d_mask(offs_m, offs_n, max_0=num_pid_m, max_1=num_pid_n)

# Calculate memory offsets for swizzled ordering
offs_sw_m = get_1d_offset(1, n_prev_chunks=pid_m_)
offs_sw_n = get_1d_offset(1, n_prev_chunks=pid_n_)
offs_sw = get_2d_offset(offs_sw_m, offs_sw_n, stride_0=num_pid_n)
mask_sw = get_2d_mask(offs_sw_m, offs_sw_n, max_0=num_pid_m, max_1=num_pid_n)

# Load from original pattern and store in swizzled pattern
x = tl.load(x_ptr + offs, mask=mask)
tl.store(z_ptr + offs_sw, x, mask=mask_sw)

```

```

# Demonstrate swizzling effect
blocks_m, blocks_n = 5, 4
x = torch.arange(blocks_m*blocks_n, device='cuda').view(blocks_m, blocks_n)
print("\nOriginal matrix:")
print(x)

```

```

tensor([[ 0,  1,  2,  3],
        [ 4,  5,  6,  7],
        [ 8,  9, 10, 11],
        [12, 13, 14, 15],
        [16, 17, 18, 19]], device='cuda:0')

```

```

z = -torch.ones_like(x) # Initialize output matrix with -1
print("\nEmpty output matrix:")
print(z)

```

```

tensor([[ -1,  -1,  -1,  -1],
        [ -1,  -1,  -1,  -1],
        [ -1,  -1,  -1,  -1],
        [ -1,  -1,  -1,  -1],
        [ -1,  -1,  -1,  -1]], device='cuda:0')

```

```

# Apply swizzling
swizzle_k[(blocks_m,blocks_n)](x, z, group_sz=3)
print("\nSwizzled matrix:")
print(z)

```

```

tensor([[ 0,  3,  6,  9],
        [ 1,  4,  7, 10],
        [ 2,  5,  8, 11],
        [12, 14, 16, 18],
        [13, 15, 17, 19]], device='cuda:0')

```

- Use `tl.swizzle2d(pid, num_pid_m, num_pid_n, group_size)`

Adjusted Kernel

- **Modify PID:**
 - Apply swizzling to **PID** before computing offsets.
- **Benefits:**
 - Reduces the number of unique data loads.
 - Increases cache hits, improving performance.
- **Grouped Matrix Multiplication with Swizzling:**

```
@triton.jit
def grouped_matmul_k(
    a_ptr, b_ptr, c_ptr,
    m, n, k,
    stride_am, stride_ak,
    stride_bk, stride_bn,
    stride_cm, stride_cn,
    bm: tl.constexpr, bn: tl.constexpr, bk: tl.constexpr,
    group_sz: tl.constexpr
):
    """
    Matrix multiplication kernel with memory access pattern optimization using swizzling
    This implementation groups thread blocks to improve cache utilization.

    Args:
        a_ptr, b_ptr: Input matrix pointers
        c_ptr: Output matrix pointer
        m, n, k: Matrix dimensions
        stride_*: Memory strides for each matrix
        bm, bn, bk: Block sizes for tiled computation
        group_sz: Size of thread block groups for swizzling
    """
    # Get thread block coordinates and grid dimensions
    pid_m, pid_n = tl.program_id(0), tl.program_id(1)
    num_pid_m, num_pid_n = tl.num_programs(0), tl.num_programs(1)

    # Apply swizzling to optimize memory access pattern
    pid_m, pid_n = tl.swizzle2d(pid_m, pid_n, num_pid_m, num_pid_n, group_sz)

    # Calculate block offsets
    rm = get_1d_offset(size=bm, n_prev_chunks=pid_m)
    rn = get_1d_offset(size=bn, n_prev_chunks=pid_n)
    rk = get_1d_offset(size=bk, n_prev_chunks=0)

    # Calculate memory offsets for input matrices
    offs_a = a_ptr + get_2d_offset(rm, rk, stride_am, stride_ak)
    offs_b = b_ptr + get_2d_offset(rk, rn, stride_bk, stride_bn)

    # Initialize accumulator
```

```

acc = tl.zeros((bm, bn), dtype=tl.float32)

# Main computation loop
for _ in range(0, k, bk):
    a = tl.load(offfs_a)
    b = tl.load(offfs_b)
    acc += tl.dot(a, b, allow_tf32=False)
    offfs_a += bk * stride_ak
    offfs_b += bk * stride_bk

# Store results
c = c_ptr + get_2d_offset(rm, rn, stride_cm, stride_cn)
mask = get_2d_mask(rm, rn, m, n)
tl.store(c, acc, mask=mask)

```

Validation

- Testing:

```

# Create simplified interface for grouped matrix multiplication
grouped_matmul = partial(matmul, matmul_k_fn=grouped_matmul_k)

```

```

# Small example
print("\nTesting with small matrices:")
a = torch.ones((3, 4), dtype=torch.float32, device='cuda')
b = torch.ones((4, 5), dtype=torch.float32, device='cuda')
grouped_matmul(a, b, group_sz=4)

```

```

tensor([[4., 4., 4., 4., 4.],
        [4., 4., 4., 4., 4.],
        [4., 4., 4., 4., 4.]], device='cuda:0', dtype=torch.float16)

```

```

# Larger example with verification
print("\nTesting with larger matrices:")
torch.manual_seed(0)
a = torch.randn((512, 512), device='cuda', dtype=torch.float16)
b = torch.randn((512, 512), device='cuda', dtype=torch.float16)

triton_output = grouped_matmul(a, b, group_sz=32)
torch_output = torch.matmul(a, b)

# Verify results
if torch.allclose(triton_output, torch_output, atol=5e-2, rtol=0):
    print("✅ Triton and Torch match")
else:
    print("❌ Triton and Torch differ")

```

✅ Triton and Torch match

- Compare output with PyTorch's `torch.matmul`.
- Use various matrix sizes for thorough testing.

Benchmarking

Purpose

- Measure and compare kernel performance.
- Identify performance gains or bottlenecks.

Tools and Methods

- Triton's Benchmarking Utilities:
 - Provides functions to benchmark kernels over input ranges.
- Parameters:
 - Test different matrix sizes and block sizes.

Findings (RTX 4090)

Benchmark #1

```
"""
Performance Benchmarking for Matrix Multiplication Implementations

Compares the performance of the three matrix multiplication implementations:
1. Naive Triton implementation
2. Grouped Triton implementation (with memory access optimization)
3. PyTorch's native implementation

The benchmark measures performance in GB/s (gigabytes per second) across different matrix
"""

@triton.testing.perf_report(
    triton.testing.Benchmark(
        # X-axis configuration
        x_names=['square_matrix_size'], # What varies along x-axis
        x_vals=[2**i for i in range(5, 12, 1)], # Matrix sizes: 32, 64, 128, 256, 512, 1024
        x_log=True, # Use logarithmic scale for x-axis

        # Different implementations to compare (creates different lines on plot)
        line_arg='provider', # Parameter that determines which implementation to use
        line_vals=['naive', 'grouped', 'torch'], # Possible implementation values
        line_names=['Naive', 'Grouped', 'Torch'], # Labels for each implementation

        # Plot styling
        styles=[('blue', '-'), ('green', '-'), ('orange', '-')], # Colors and line styles
        ylabel='GB/s', # Y-axis label showing throughput
```

```

    plot_name='matmul-performance', # Name for saving the plot

    # Additional arguments (empty in this case)
    args={},
))
def benchmark(square_matrix_size, provider):
    """
    Benchmark different matrix multiplication implementations.

    Args:
        square_matrix_size: Size of the square matrices to multiply (N×N)
        provider: Which implementation to benchmark ('naive', 'grouped', or 'torch')

    Returns:
        tuple: (median_performance, min_performance, max_performance) in GB/s

    Performance calculation:
    - Matrix multiplication requires reading 2 matrices and writing 1 matrix
    - Each matrix has size N×N with 4 bytes per element (float32)
    - Total memory moved = 3 * N * N * 4 bytes
    - GB/s = (12 * N * N) / (time_in_ms * 1e6) # 12 = 3 matrices * 4 bytes
    """
    # Create random input matrices
    sz = square_matrix_size
    a = torch.rand((sz, sz), device='cuda', dtype=torch.float32)
    b = torch.rand((sz, sz), device='cuda', dtype=torch.float32)

    # Define measurement percentiles
    quantiles = [0.5, 0.2, 0.8] # median, 20th percentile, 80th percentile


    # Benchmark the requested implementation
    if provider == 'naive':
        ms, min_ms, max_ms = triton.testing.do_bench(
            lambda: naive_matmul(a, b),
            quantiles=quantiles
        )
    if provider == 'grouped':
        ms, min_ms, max_ms = triton.testing.do_bench(
            lambda: grouped_matmul(a, b, group_sz=8),
            quantiles=quantiles
        )
    if provider == 'torch':
        ms, min_ms, max_ms = triton.testing.do_bench(
            lambda: torch.matmul(a,b),
            quantiles=quantiles
        )

    # Convert milliseconds to GB/s
    gbps = lambda ms: 12 * sz * sz / ms * 1e-6 # Formula explained in docstring

```

```
return gbps(ms), gbps(max_ms), gbps(min_ms)
```

```
# Run the benchmark
print("\nRunning performance benchmark...")
print("This will test matrix sizes from 32×32 to 2048×2048")
print("For each size, we'll compare naive Triton, grouped Triton, and PyTorch implementat")
benchmark.run(print_data=True, show_plots=True)
```



matmul-performance:

	square_matrix_size	Naive	Grouped	Torch
0	32.0	0.093750	0.093750	0.046875
1	64.0	0.187500	0.158940	0.083333
2	128.0	0.290909	0.284024	0.214286
3	256.0	0.333333	0.333333	0.300000
4	512.0	0.190855	0.193548	0.300000
5	1024.0	0.055379	0.055339	0.218182
6	2048.0	0.014159	0.014179	0.079470

Benchmark #2

```
# ===== Impact of Batch Size =====
@triton.testing.perf_report(
    triton.testing.Benchmark(
        # X-axis: varying batch sizes
```

```

x_names=['batch_size'],
x_vals=[2**i for i in range(4, 7, 1)], # Testing batch sizes: 16, 32, 64
x_log=True,

# Compare different implementations
line_arg='provider',
line_vals=['naive', 'grouped', 'torch'],
line_names=['Naive', 'Grouped', 'Torch'],

# Plot styling
styles=[('blue', '-'), ('green', '-'), ('orange', '-')],
ylabel='GB/s',
plot_name='matmul-performance',
args={}
))
def benchmark(batch_size, provider):
    """
    Benchmark matrix multiplication with varying batch sizes.

    This benchmark keeps matrix size fixed (512x512) and varies the computation
    batch size to understand its impact on performance.

    Args:
        batch_size: Size of computation batches (block size)
        provider: Which implementation to benchmark

    Returns:
        tuple: (median_performance, min_performance, max_performance) in GB/s
    """
    # Fixed matrix size of 512x512
    sz = 512
    a = torch.rand((sz, sz), device='cuda', dtype=torch.float32)
    b = torch.rand((sz, sz), device='cuda', dtype=torch.float32)

    quantiles = [0.5, 0.2, 0.8]


    # Benchmark each implementation with varying batch sizes
    if provider == 'naive':
        ms, min_ms, max_ms = triton.testing.do_bench(
            lambda: naive_matmul(a, b, bs=batch_size),
            quantiles=quantiles
        )
    if provider == 'grouped':
        ms, min_ms, max_ms = triton.testing.do_bench(
            lambda: grouped_matmul(a, b, bs=batch_size, group_sz=8),
            quantiles=quantiles
        )
    if provider == 'torch':
        ms, min_ms, max_ms = triton.testing.do_bench(
            lambda: torch.matmul(a,b),
            quantiles=quantiles

```

```
)
```

```
gbps = lambda ms: 12 * sz / ms * 1e-6  
return gbps(ms), gbps(max_ms), gbps(min_ms)
```

```
# Run the benchmark  
print("\nRunning batch size impact benchmark...")  
print("Testing different batch sizes on 512x512 matrices")  
benchmark.run(print_data=True, show_plots=True)
```



matmul-performance:

	batch_size	Naive	Grouped	Torch
0	16.0	0.171429	0.176471	0.3
1	32.0	0.375000	0.360902	0.3
2	64.0	0.352941	0.352941	0.3

Benchmark #3

```
# ===== Matrix Size Impact with Fixed Block Size =====
```

```
@triton.testing.perf_report(  
    triton.testing.Benchmark(  
        # X-axis: varying matrix sizes  
        x_names=['square_matrix_size'],  
        x_vals=[2**i for i in range(5, 12, 1)], # Matrix sizes from 32 to 2048
```



```

x_log=True,

# Compare different implementations
line_arg='provider',
line_vals=['naive', 'grouped', 'torch'],
line_names=['Naive', 'Grouped', 'Torch'],

# Plot styling
styles=[('blue', '-'), ('green', '-'), ('orange', '-')],
ylabel='GB/s',
plot_name='matmul-performance',
args={}
))
def benchmark(square_matrix_size, provider):
    """
    Benchmark matrix multiplication with varying matrix sizes but fixed block size.

    This benchmark uses a fixed block size (64) while varying matrix dimensions
    to understand how different implementations scale with problem size.

    Args:
        square_matrix_size: Size of the square matrices to multiply (N×N)
        provider: Which implementation to benchmark

    Returns:
        tuple: (median_performance, min_performance, max_performance) in GB/s
    """
    sz = square_matrix_size
    a = torch.rand((sz, sz), device='cuda', dtype=torch.float32)
    b = torch.rand((sz, sz), device='cuda', dtype=torch.float32)

    quantiles = [0.5, 0.2, 0.8]

    # Fixed block size of 64 for all implementations
    if provider == 'naive':
        ms, min_ms, max_ms = triton.testing.do_bench(
            lambda: naive_matmul(a, b, bs=64),
            quantiles=quantiles
        )
    if provider == 'grouped':
        ms, min_ms, max_ms = triton.testing.do_bench(
            lambda: grouped_matmul(a, b, group_sz=8, bs=64),
            quantiles=quantiles
        )
    if provider == 'torch':
        ms, min_ms, max_ms = triton.testing.do_bench(
            lambda: torch.matmul(a,b),
            quantiles=quantiles
        )

```

```

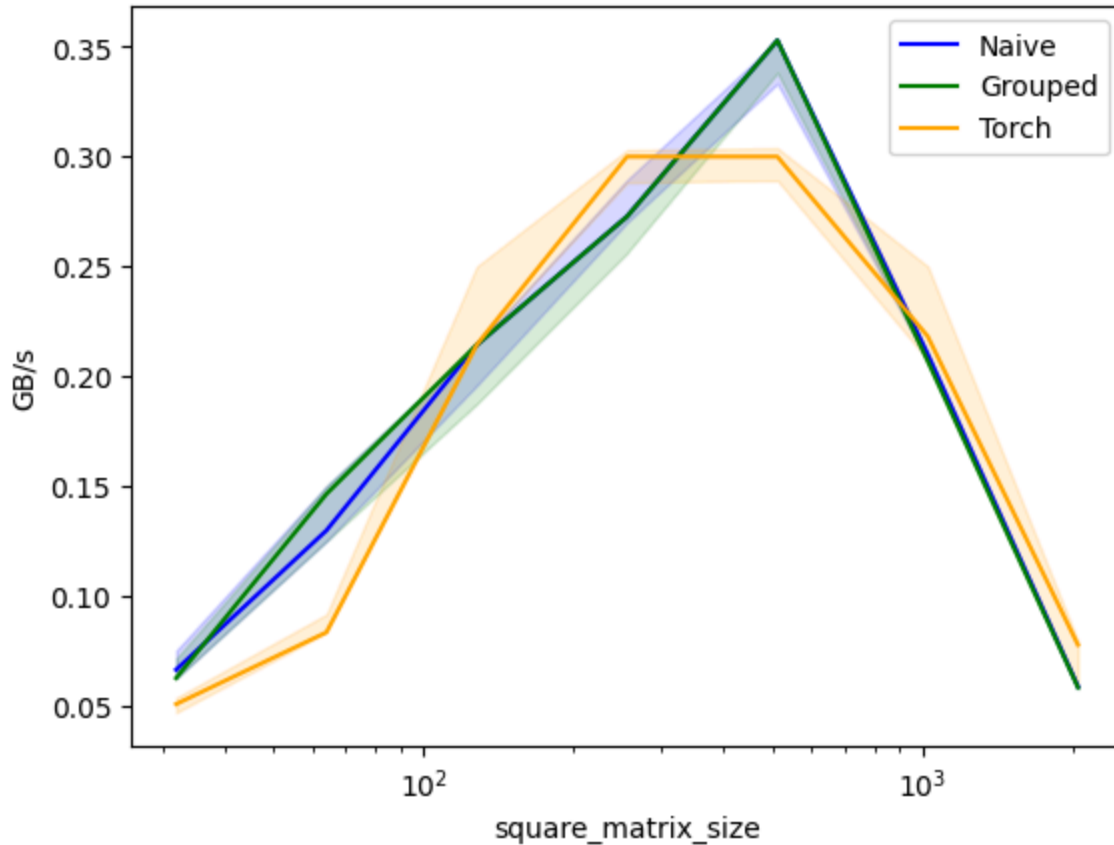
gbps = lambda ms: 12 * sz / ms * 1e-6
return gbps(ms), gbps(max_ms), gbps(min_ms)

```

```

print("\nRunning matrix size scaling benchmark...")
print("Testing different matrix sizes with fixed block size=64")
benchmark.run(print_data=True, show_plots=True)

```



matmul-performance:

	square_matrix_size	Naive	Grouped	Torch
0	32.0	0.066298	0.062500	0.050633
1	64.0	0.129730	0.146341	0.083333
2	128.0	0.214286	0.214286	0.214286
3	256.0	0.272727	0.272727	0.300000
4	512.0	0.352941	0.352941	0.300000
5	1024.0	0.210066	0.206897	0.218182
6	2048.0	0.058680	0.058252	0.077623

- **Performance Trends:**
 - **Small Matrices:**
 - Triton kernels can outperform PyTorch.
 - **Large Matrices:**
 - PyTorch may be faster due to highly optimized kernels.
- **Cache Effects:**
 - Performance drops when exceeding **L1** or **L2** cache capacity.

- **Block Size Impact:**
 - Larger block sizes generally improve performance.
 - Excessively large block sizes may cause **out-of-memory** errors.

Profiling Tools

- **NVIDIA Nsight Compute (NCU):**
 - Provides detailed performance metrics.
 - Helps identify optimization opportunities.

Auto-Tuning

Concept

- **Auto-Tuning:**
 - Automatically finds the best kernel configurations for performance.
- **Parameters Tuned:**
 - **Block sizes, tile sizes, and other kernel parameters.**

Implementation in Triton

```
"""
```

Matrix Multiplication with Autotuning

This implementation adds automatic performance tuning by testing different configurations of block sizes, group sizes, and other parameters to find optimal settings for different matrix sizes.

```
"""
```

```
# ===== Autotuned Matrix Multiplication Kernel =====
```

```
@triton.autotune(
```

```
    configs=[
```

```
        # Different configurations to try, varying block sizes and execution parameters
        triton.Config({'bm': 128, 'bn': 256, 'bk': 64, 'group_sz': 8},
                       num_stages=3, num_warps=8), # Larger blocks, fewer stages
        triton.Config({'bm': 64, 'bn': 256, 'bk': 32, 'group_sz': 8},
                       num_stages=4, num_warps=4),
        triton.Config({'bm': 128, 'bn': 128, 'bk': 32, 'group_sz': 8},
                       num_stages=4, num_warps=4),
        triton.Config({'bm': 128, 'bn': 64, 'bk': 32, 'group_sz': 8},
                       num_stages=4, num_warps=4),
        triton.Config({'bm': 64, 'bn': 128, 'bk': 32, 'group_sz': 8},
                       num_stages=4, num_warps=4),
        triton.Config({'bm': 128, 'bn': 32, 'bk': 32, 'group_sz': 8},
                       num_stages=4, num_warps=4),
        triton.Config({'bm': 64, 'bn': 32, 'bk': 32, 'group_sz': 8},
                       num_stages=5, num_warps=2), # Smaller blocks, more stages
```

```

        triton.Config({'bm': 32, 'bn': 64, 'bk': 32, 'group_sz': 8},
                      num_stages=5, num_warps=2),
    ],
    # Autotuning is based on input matrix dimensions
    key=['m', 'n', 'k'],
)
@triton.jit
def grouped_autotuned_matmul_k(
    a_ptr, b_ptr, c_ptr,
    m, n, k,
    stride_am, stride_ak,
    stride_bk, stride_bn,
    stride_cm, stride_cn,
    bm: tl.constexpr, bn: tl.constexpr, bk: tl.constexpr,
    group_sz: tl.constexpr
):
    """
    Autotuned matrix multiplication kernel that tries different configurations
    to find the best performance for given matrix dimensions.

    The configurations vary:
    - Block sizes (bm, bn, bk)
    - Number of pipeline stages
    - Number of warps
    - Group size for memory access optimization
    """
    # Get thread block coordinates and grid dimensions
    pid_m = tl.program_id(0)
    pid_n = tl.program_id(1)
    num_pid_m = tl.num_programs(0)
    num_pid_n = tl.num_programs(1)

    # Apply swizzling for memory access optimization
    pid_m, pid_n = tl.swizzle2d(pid_m, pid_n, num_pid_m, num_pid_n, group_sz)

    # Calculate block offsets
    rm = get_1d_offset(size=bm, n_prev_chunks=pid_m)
    rn = get_1d_offset(size=bn, n_prev_chunks=pid_n)
    rk = get_1d_offset(size=bk, n_prev_chunks=0)

    # Calculate memory offsets
    offs_a = a_ptr + get_2d_offset(rm, rk, stride_am, stride_ak)
    offs_b = b_ptr + get_2d_offset(rk, rn, stride_bk, stride_bn)

    # Matrix multiplication computation
    acc = tl.zeros((bm, bn), dtype=tl.float32)
    for _ in range(0, k, bk):
        a = tl.load(offs_a)
        b = tl.load(offs_b)
        acc += tl.dot(a, b, allow_tf32=False)
        offs_a += bk * stride_ak

```

```

        offs_b += bk * stride_bk

# Store results
c = c_ptr + get_2d_offset(rm, rn, stride_cm, stride_cn)
mask = get_2d_mask(rm, rn, m, n)
tl.store(c, acc, mask=mask)

```

```

def grouped_autotuned_matmul(a, b):
    """
    High-level wrapper for autotuned matrix multiplication.

    This function handles:
    1. Input validation
    2. Output initialization
    3. Grid computation
    4. Kernel launch with autotuned parameters
    """
    matmul_k_fn = grouped_autotuned_matmul_k

    # Validate inputs
    assert a.shape[1] == b.shape[0], "matrix dims not compatible for matmul"
    check_tensors_gpu_ready(a, b)

    # Get matrix dimensions
    (m, k), (_, n) = a.shape, b.shape

    # Initialize output matrix
    c = torch.empty((m, n), device=a.device, dtype=torch.float16)

    # Compute grid dimensions
    grid = lambda meta: (triton.cdiv(m, meta['bm']), triton.cdiv(n, meta['bn']))

    # Launch kernel with autotuned parameters
    matmul_k_fn[grid](
        a, b, c,
        m, n, k,
        a.stride(0), a.stride(1),
        b.stride(0), b.stride(1),
        c.stride(0), c.stride(1),
        # Block sizes and group size are autotuned
    )
    return c

```

```

a,b = torch.ones(3,4, device='cuda'), torch.ones(4,5, device='cuda')
a@b

```

```

tensor([[4., 4., 4., 4., 4.],
        [4., 4., 4., 4., 4.],
        [4., 4., 4., 4., 4.]], device='cuda:0')

```

```
grouped_autotuned_matmul(a,b)
```

```
tensor([[4., 4., 4., 4., 4.],
        [4., 4., 4., 4., 4.],
        [4., 4., 4., 4., 4.]], device='cuda:0', dtype=torch.float16)
```

- **Define Configurations:**
 - List possible values for parameters.
- **Auto-Tuner Decorator:**
 - Use `@triton.autotune(configs=..., key=['M', 'N', 'K'])`.
- **Execution:**
 - Triton tests each configuration to find the optimal one.
 - The best configuration is cached for reuse.

Observations

```
# ===== Performance Benchmark with Autotuning =====

@triton.testing.perf_report(
    triton.testing.Benchmark(
        x_names=['square_matrix_size'],
        x_vals=[2**i for i in range(5, 12, 1)], # 32 to 2048
        x_log=True,
        line_arg='provider',
        line_vals=['naive', 'grouped', 'grouped-autotuned', 'torch'],
        line_names=['Naive', 'Grouped', 'Grouped & Auto-Tuned', 'Torch'],
        styles=[('blue', '-'), ('green', '-'), ('green', '--'), ('orange', '-')],
        ylabel='GB/s',
        plot_name='matmul-performance',
        args={}
    )
)

def benchmark(square_matrix_size, provider):
    """
    Benchmark comparing all implementations including autotuned version.

    Compares:
    1. Naive Triton implementation
    2. Grouped Triton implementation
    3. Grouped & Autotuned Triton implementation
    4. PyTorch native implementation
    """
    # Create test matrices
    sz = square_matrix_size
    a = torch.rand((sz, sz), device='cuda', dtype=torch.float32)
    b = torch.rand((sz, sz), device='cuda', dtype=torch.float32)

    quantiles = [0.5, 0.2, 0.8]
```

```

# Benchmark each implementation
if provider == 'naive':
    ms, min_ms, max_ms = triton.testing.do_bench(
        lambda: naive_matmul(a, b, bs=64),
        quantiles=quantiles
    )
if provider == 'grouped':
    ms, min_ms, max_ms = triton.testing.do_bench(
        lambda: grouped_matmul(a, b, group_sz=8, bs=64),
        quantiles=quantiles
    )
if provider == 'grouped-autotuned':
    ms, min_ms, max_ms = triton.testing.do_bench(
        lambda: grouped_autotuned_matmul(a, b),
        quantiles=quantiles
    )
if provider == 'torch':
    ms, min_ms, max_ms = triton.testing.do_bench(
        lambda: torch.matmul(a,b),
        quantiles=quantiles
    )

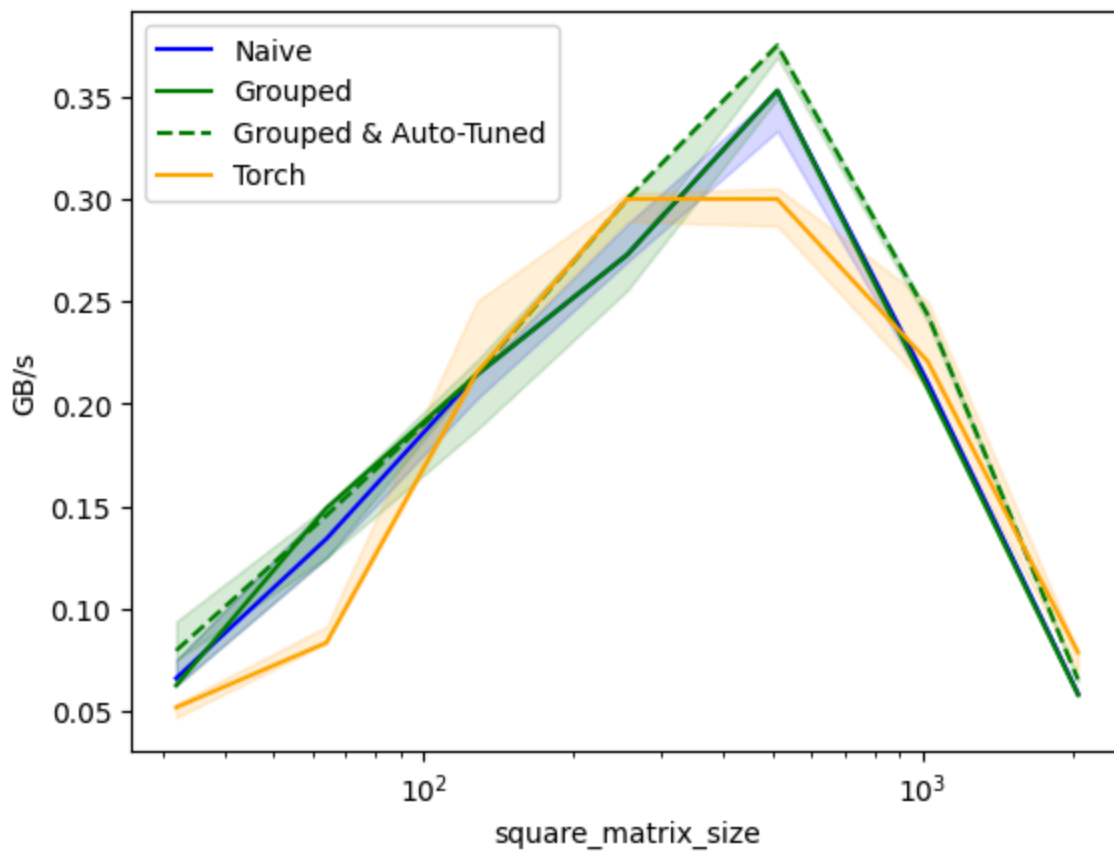
gbps = lambda ms: 12 * sz / ms * 1e-6
return gbps(ms), gbps(max_ms), gbps(min_ms)

```

```

# Run the benchmark
print("\nRunning final performance comparison with autotuning...")
benchmark.run(print_data=True, show_plots=True)

```



matmul-performance:

	square_matrix_size	Naive	Grouped	Grouped & Auto-Tuned	Torch
0	32.0	0.065934	0.062500	0.079470	0.051724
1	64.0	0.134078	0.149068	0.145455	0.083333
2	128.0	0.214286	0.214286	0.214286	0.215247
3	256.0	0.272727	0.272727	0.300000	0.300000
4	512.0	0.352941	0.352941	0.375000	0.300000
5	1024.0	0.210526	0.207343	0.243500	0.220753
6	2048.0	0.058492	0.057971	0.065362	0.078689

- **Performance Improvements:**
 - Auto-tuning can significantly enhance performance.
- **Unexpected Results:**
 - In some cases, auto-tuned kernels may perform worse.
 - Requires analysis to adjust configurations.

Tips

- **Problem Size Specificity:**
 - Optimal configurations may vary with input sizes.
- **Best Practices:**
 - Refer to Triton documentation and community resources.
 - Experiment with different configurations.

Conclusion and Resources

Summary

- Triton provides an accessible way to write efficient GPU kernels.
- Offers a balance between **ease of use** and **performance**.
- **Debugging** and **auto-tuning** tools enhance development.
- **Benchmarking** is essential for performance validation.

Further Learning

- **Resources:**
 - [Triton Documentation](#): Comprehensive guide and reference.
 - **Lectures and Talks:**
 - [Lecture 1: How to profile CUDA kernels in PyTorch](#)
 - [Notes](#)
 - [Lecture 9: Reductions](#)
 - [Notes](#)
 - **LightLLM Triton Kernels:** [lightllm/common/basemodel/triton_kernel](#)
 - **unsloth Triton Kernels:** [unsloth/kernels](#)
 - **Triton Puzzles:** [srush/Triton-Puzzles](#)

About Me:

I'm Christian Mills, a deep learning consultant specializing in practical AI implementations. I help clients leverage cutting-edge AI technologies to solve real-world problems.

Interested in working together? Fill out my [Quick AI Project Assessment](#) form or learn more [about me](#).

0 Comments - powered by [utteranc.es](#)